

This document outlines the kube yaml fields that are currently supported by the **podman kube play** command. The notes were cross referenced with the Kubernetes Docs, particularly the [Greppable Reference](#), where printable view make Ctrl-F go vroom.

Not Supported ❌ Not Documented ? Supported ✅ Only with multiple nodes ☑

1. Pod Fields

PodSpec

Field		Field	
containers	✅	os.name	❌
initContainers	✅	volumes	❌
imagePullSecrets	❌		
enableServiceLinks	❌		

1.1. Scheduling

[KubeSchedulerConfiguration](#): has many [profiles](#), each identified by a `schedulerName`. See also section on [Scheduling Profiles](#).

Field		Field		Field	
nodeSelector	☑	priority	❌	preemptionPolicy	?
nodeName	☑	priorityClassName	❌	overhead	?
schedulerName	☑	runtimeClassName	❌		

- [PriorityClass](#)
- [RuntimeClass](#)

1.1.1. Affinity, Tolerations, Taints, Topology Spread Constraints

affinity.*		tolerations.*		topologySpreadConstraints.*	
nodeAffinity	☑	key	☑	maxSkew	☑
podAffinity	☑	operator	☑	topologyKey	☑
podAntiAffinity	☑	effect	☑	whenUnsatisfiable	☑
		tolerationSeconds	☑	labelSelector	☑
			☑	minDomains	☑

1.1.2. Lifecycle

Field	
restartPolicy	✅
terminationGracePeriod	❌
activeDeadlineSeconds	❌
readinessGates.conditionType	❌

1.2. DNS

Field		hostAliases.*		dnsConfig.*	
hostname	✅	hostAliases.hostnames	✅	nameservers	✅
setHostnameAsFQDN	❌	hostAliases.ip	✅	options.name	✅
subdomain	❌			options.value	✅
dnsPolicy	❌			searches	✅

**hostAliases**  
only valid for non-hostNetwork pods.

1.3. Security

Service Accounts

Field	
serviceAccountName	❌
automountServiceAccountToken	❌

1.4. Namespaces, sysctl and seccomp profiles

[Emacs also supports seccomp profiles](#), if you're paranoid and for some reason compiling BPF programs isn't an undue time investment.

Field		securityContext.*	
hostNetwork	✅	seccompProfile.type	❌
hostPID	✅	seccompProfile.localhostProfile	❌
hostIPC	✅	sysctls.name	❌
shareProcessNamespace	✅	sysctls.value	❌

1.5. Volume/Process Ownership and SELinux

Other `securityContext.*` fields for PodSpec

securityContext.*		securityContext.*	
runAsUser	✅	seLinuxOptions.level	✅
runAsNonRoot	❌	seLinuxOptions.role	✅
runAsGroup	✅	seLinuxOptions.type	✅
supplementalGroups	✅	seLinuxOptions.user	✅
fsGroup	❌	windowsOptions.gmsaCredentialSpec	❌
fsGroupChangePolicy	❌	windowsOptions.hostProcess	❌
		windowsOptions.runAsUserName	❌

**fsGroup/ChangePolicy**  
these pertain to whether a pod can change the ownership of volumes before "being exposed inside Pod." The former is a Group ID and `setgid` will be set, changing ownership of files created. The latter can only be set to `OnRootMismatch` Or `Always`.  
Not entirely sure, but I wouldn't want to have to fix whatever this solves.

2. Container Fields

Name/Image		Entry Point		Ports		Debugging	
name	✅	command	✅	ports.containerPort	✅	stdin	❌
image	✅	args	✅	ports.hostIP	✅	stdinOnce	❌
imagePullPolicy	✅	workingDir	✅	ports.hostPort	✅	tty	❌
				ports.name	✅		
				ports.protocol	✅		

2.1. Env References

Either set `env.value` or supply a reference source with `env.valueFrom`, which needs a corresponding `envFrom.*` source

env.*		env.valueFrom.*		env.valueFrom.*		env.valueFrom.*	
name	✅	fieldRef	✅	configMapKeyRef.key	✅	secretKeyRef.key	✅
value	✅	resourceFieldRef	✅	configMapKeyRef.name	✅	secretKeyRef.name	✅
				configMapKeyRef.optional	✅	secretKeyRef.optional	✅

**env.valueFrom.fieldRef**

- Composed of fieldPath & apiVersion ([ObjectFieldSelector](#))
- Selects a field of the pod
- Only annotations, labels, name and namespace are supported.

env.valueFrom.resourceFieldRef

- Composed of resource, containerName, divisor ([ResourceFieldSelector](#))
- Selects a resource of the container
- Only resources limits and requests (limits.cpu, limits.memory, requests.cpu and requests.memory) are currently supported.

2.2. Env Sources

envFrom.*	envFrom.*	env.*
prefix	✗ configMapRef.name	✔ secretRef.name
	configMapRef.optional	✔ secretRef.optional

2.3. Volumes, Devices and Resources

volumeMounts.*	volumeMounts.*	volumeMounts.*	volumeDevices.*	resources.*
name	✔ mountPath	✔ subPath	✗ devicePath	✗ limits
readOnly	✔ mountPropagation	✗ subPathExpr	✗ name	✗ requests

2.4. Lifecycle, Termination and Probes

<a href="#">LifecycleHandler</a>	TerminationMessage	Probe
lifecycle.postStart	✗ terminationMessagePath	✗ livenessProbe
lifecycle.preStop	✗ terminationMessagePolicy	✗ readinessProbe
		✗ startupProbe

For the termination message, the **path** is the mounted file to which the container’s termination message will be written and the **policy** indicates how the termination message should be populated.

securityContext.*	securityContext.*
securityContext.runAsUser	✔ securityContext.readOnlyRootFilesystem
securityContext.runAsNonRoot	✗ securityContext.procMount
securityContext.runAsGroup	✔ securityContext.privileged
	securityContext.allowPrivilegeEscalation

2.4.1. Capabilities and Seccomp

securityContext.\*

sc.capabilities.*	sc.seccompProfile.*
add	✔ type
drop	✔ localhostProfile

2.4.2. SELinux and Windows

securityContext.\*

sc.seLinuxOptions.*	sc.windowsOptions.*
level	✔ gmsaCredentialSpec
role	✔ hostProcess
type	✔ runAsUserName
user	✔

3. PersistentVolumeClaim Fields

Field	resources.*
volumeName	✗ limits
storageClassName	✔ requests
volumeMode	✗ claims.name
accessModes	✔
selector	✗

4. ConfigMap Fields

Field
binaryData
data
immutable

5. Deployment Fields

Field	Field
replicas	✔ minReadySeconds
selector	✔ progressDeadlineSeconds
template	✔ strategy.type
revisionHistoryLimit	✗ strategy.rollingUpdate.maxSurge
paused	✗ strategy.rollingUpdate.maxUnavailable

- (Podman) For replicas the actual replica count is ignored and set to 1

Created: 2023-07-14 Fri 05:07

[Validate](#)